

NAME: MOHAMMED
MUSHARRF
SRN:
PES2UG23CS915

ML_LAB6_SEC_F 

Description

1. Introduction

The aim of this lab was to design and implement an (ANN) entirely from the ground up, without using external deep learning libraries like or . The focus was on building the fundamental components—forward propagation, backpropagation, weight initialization, and gradient descent training—manually. Key goals included:

- Implementing ReLU and Sigmoid activation functions
- Using for weight setup
- Applying (MSE) as the loss function
- Training the network using gradient descent from scratch

Part A concentrated on creating a baseline ANN, while Part B studied how changing hyperparameters such as learning rate, batch size, epochs, and activation function affected the model's accuracy and convergence.

2. Dataset Description

A synthetic dataset was generated using the student SRN as a random seed. The target function combined cubic and inverse terms to create a challenging regression problem.

- **Total Samples:** 100,000
 - **Split:** 80,000 for training and 20,000 for testing
 - **Noise:** Gaussian noise with variance 2.49
 - **Preprocessing:** Both input and output values were standardized using normalization (via)
-

3. Methodology

The neural network architecture used was **Input(1) → Hidden(72) → Hidden(32) → Output(1)**.

Implementation Details:

- **Initialization:** Weights initialized using Xavier method; biases set to zero.
 - **Forward Pass:** Hidden layers used either ReLU or Sigmoid activation, while the output layer was linear.
 - **Loss Function:** MSE
 - **Optimization:** Backpropagation implemented manually with gradient descent and optional early stopping
 - **Experiments (Part B):** Hyperparameters (learning rate, batch size, epochs, activation function) were varied systematically to observe their effect on training.
-

4. Results and Analysis

Baseline Model:

```
=====
Final Training Loss: 0.298561
Final Test Loss: 0.305123
R2 Score: 0.6997
```

Observations from Experiments:

- **Higher Learning Rate (Exp1):** Led to faster convergence and best overall accuracy ($R^2 \approx 0.87$).
- **Larger Batch Size (Exp2):** Did not produce any noticeable improvement; performance slightly dropped.
- **More Epochs (Exp3):** Training longer reduced both train/test losses and improved R^2 compared to baseline.
- **Sigmoid Activation (Exp4):** Performed very poorly due to vanishing gradients, showing a major drop in accuracy.

Results Table

===== RESULTS TABLE =====							
	Experiment	Learning Rate	Batch Size	Epochs	Activation	Final Train Loss	\
0	Exp1	0.005	32	500	relu	0.129063	
1	Exp2	0.001	64	500	relu	0.291003	
2	Exp3	0.001	32	1000	relu	0.189895	
3	Exp4	0.001	32	500	sigmoid	0.957797	
	Final Test Loss	R2 Score					
0	0.130640	0.871431					
1	0.297730	0.706991					
2	0.193456	0.809612					
3	0.973441	0.041994					

5. Conclusion

This lab showed how an ANN can be fully developed from scratch to solve a regression problem. The baseline model gave good performance on the synthetic dataset, and hyperparameter tuning revealed:

- Learning rate and number of epochs are critical for achieving better convergence and accuracy.
- ReLU activation works significantly better than Sigmoid in deeper layers, especially for this regression task.

Overall, the experiments highlighted how much model effectiveness depends on choosing suitable hyperparameters and activation functions.

```
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler

##PART -A
```

AUTOMATIC ASSIGNMENT BASED ON SRN - DO NOT MODIFY

```
def get_student_assignment(student_id):
    """
    Generate unique polynomial TYPE and architecture based on student
    ID
    Uses last 3 digits of student ID for assignment
    """

    last_three = int(student_id[-3:])

    poly_type = last_three % 5

    np.random.seed(last_three)

    if poly_type == 0:
        degree = 2
    a = 0
        b = np.random.uniform(0.8, 1.5)
    c = np.random.uniform(3.0, 8.0)
    d = np.random.uniform(5.0, 15.0)
        poly_desc = f"QUADRATIC: y = {b:.2f}x² + {c:.2f}x + {d:.2f}"

    elif poly_type == 1:
        degree = 3
        a = np.random.uniform(1.8, 2.5)
    b = np.random.uniform(-1.2, 0.2)
    c = np.random.uniform(3.0, 6.0)
    d = np.random.uniform(8.0, 12.0)
        poly_desc = f"CUBIC: y = {a:.2f}x³ + {b:.2f}x² + {c:.2f}x +
{d:.2f}"

    elif poly_type == 2:
        degree = 4
        a = np.random.uniform(0.008, 0.02)
    b = np.random.uniform(1.5, 2.2)
```

```

        c = np.random.uniform(-1.0, 0.5)
d = np.random.uniform(2.0, 5.0)
e = np.random.uniform(8.0, 12.0)
    poly_desc = f"QUARTIC: y = {a:.4f}x⁴ + {b:.2f}x³ + {c:.2f}x² +
{d:.2f}x + {e:.2f}"
    coefficients = (a, b, c, d, e)

    elif poly_type == 3:
degree = "sine"
        a = np.random.uniform(1.5, 2.8)
b = np.random.uniform(-0.8, 0.8)           c =
np.random.uniform(3.0, 6.0)           d =
np.random.uniform(8.0, 12.0)          freq =
np.random.uniform(0.02, 0.05)         amp =
np.random.uniform(5.0, 15.0)
    poly_desc = f"CUBIC + SINE: y = {a:.2f}x³ + {b:.2f}x² +
{c:.2f}x + {d:.2f} + {amp:.1f}*sin({freq:.3f}x)"
coefficients = (a, b, c, d, freq, amp)

    else:           degree =
"inverse"
        a = np.random.uniform(1.8, 2.5)
b = np.random.uniform(-1.0, 0.5)           c =
np.random.uniform(3.0, 6.0)           d =
np.random.uniform(8.0, 12.0)          inv_coeff
= np.random.uniform(50, 200)
    poly_desc = f"CUBIC + INVERSE: y = {a:.2f}x³ + {b:.2f}x² +
{c:.2f}x + {d:.2f} + {inv_coeff:.1f}/x"
coefficients = (a, b, c, d, inv_coeff)

    if poly_type in [0, 1]:
coefficients = (a, b, c, d)      noise_std =
np.random.uniform(1.5, 2.5)

    arch_type = last_three % 4
architectures = {
    0: {"hidden1": 64, "hidden2": 64, "lr": 0.001, "batch_desc":
"Balanced Architecture"},
    1: {"hidden1": 32, "hidden2": 72, "lr": 0.005, "batch_desc":
"Narrow-to-Wide Architecture"},
    2: {"hidden1": 72, "hidden2": 32, "lr": 0.001, "batch_desc":
"Wide-to-Narrow Architecture"},
    3: {"hidden1": 96, "hidden2": 96, "lr": 0.003, "batch_desc":
"Large Balanced Architecture"}
}

```

```

    return {
        "polynomial_type": poly_type,
        "degree": degree,
        "coefficients": coefficients,
        "polynomial_desc": poly_desc,
        "noise_std": noise_std,
        "architecture": architectures[arch_type],
        "student_seed": last_three
    }

# Get your assignment
assignment = get_student_assignment(STUDENT_ID)
poly_type = assignment["polynomial_type"]
degree = assignment["degree"]
coefficients = assignment["coefficients"]
noise_std = assignment["noise_std"]
hidden1 = assignment["architecture"]["hidden1"]
hidden2 = assignment["architecture"]["hidden2"]
learning_rate = assignment["architecture"]["lr"]

print("=". * 70)
print(f"ASSIGNMENT FOR STUDENT ID: {STUDENT_ID}")
print("=". * 70)
print(f"Polynomial Type: {assignment['polynomial_desc']}")
print(f"Noise Level: ε ~ N(0, {noise_std:.2f})")
print(f"Architecture: Input(1) → Hidden({hidden1}) → Hidden({hidden2})"
      → Output(1)")
print(f"Learning Rate: {learning_rate}")
print(f"Architecture Type: {assignment['architecture']"
      ['batch_desc']}")
print("=". * 70)
=====
```

```
ASSIGNMENT FOR STUDENT ID: PES2UG23CS915
=====
```

```
Polynomial Type: QUARTIC:  $y = 0.0112x^4 + 1.70x^3 + 0.32x^2 + 4.27x + 10.67$ 
```

```
Noise Level: ε ~ N(0, 2.27)
```

```
Architecture: Input(1) → Hidden(96) → Hidden(96) → Output(1)
Learning Rate: 0.003
```

```
Architecture Type: Large Balanced Architecture
=====
```

DATASET GENERATION - DO NOT MODIFY

```
np.random.seed(assignment["student_seed"])

n_samples = 100000
```

```

x      = np.random.uniform(-100, 100, n_samples)

if poly_type == 0:
    _, b, c, d = coefficients
y      = b * x**2 + c * x + d + np.random.normal(0, noise_std,
n_samples)

elif poly_type == 1:      a,
b, c, d = coefficients
    y = a * x**3 + b * x**2 + c * x + d + np.random.normal(0,
noise_std, n_samples)

elif poly_type == 2:      a, b,
c, d, e = coefficients
    y = a * x**4 + b * x**3 + c * x**2 + d * x + e +
np.random.normal(0, noise_std, n_samples)

elif poly_type == 3:      a, b, c, d,
freq, amp = coefficients
    y = a * x**3 + b * x**2 + c * x + d + amp * np.sin(freq * x) +
np.random.normal(0, noise_std, n_samples)

else:      a, b, c, d, inv_coeff =
coefficients

    y = a * x**3 + b * x**2 + c * x + d + inv_coeff / (x + np.sign(x)
* 0.1) + np.random.normal(0, noise_std, n_samples)

df = pd.DataFrame({'x': x, 'y': y})
df.to_csv('student_polynomial_dataset.csv', index=False)
print(f"Dataset with {n_samples:,} samples generated and saved!")

X = df['x'].values.reshape(-1, 1)
Y = df['y'].values.reshape(-1, 1)

X_train, X_test, Y_train, Y_test = train_test_split(X, Y,
test_size=0.2, random_state=42)

scaler_X = StandardScaler()
scaler_Y = StandardScaler()

X_train_scaled = scaler_X.fit_transform(X_train)
X_test_scaled = scaler_X.transform(X_test)
Y_train_scaled = scaler_Y.fit_transform(Y_train)
Y_test_scaled = scaler_Y.transform(Y_test)

print(f"Training samples: {len(X_train_scaled):,}")
print(f"Test samples: {len(X_test_scaled):,}")

```

```
Dataset with 100,000 samples generated and saved!
Training samples: 80,000
Test samples: 20,000
```

ACTIVATION FUNCTIONS- TODO: IMPLEMENT

```
# ACTIVATION FUNCTIONS
def relu(z):
    """ReLU activation"""
    return np.maximum(0, z)
def
relu_derivative(z):
    """Derivative of ReLU"""
    return (z > 0).astype(float)
```

LOSS FUNCTION- TODO: IMPLEMENT

```
def mse_loss(y_true, y_pred):
    """Mean Squared Error"""
    return np.mean((y_true - y_pred) ** 2)
```

WEIGHT INITIALIZATION - TODO: IMPLEMENT

XAVIER INITIALIZATION

Xavier (Glorot) Initialization

When training neural networks, how we initialize weights matters.

- If weights are **too small** → activations and gradients vanish.
- If weights are **too large** → activations and gradients explode.

#Xavier initialization (Glorot & Bengio, 2010) balances this by keeping the variance of activations roughly the same across all layers.

Formula

Let:

- **fan_in** = number of input units to a layer
- **fan_out** = number of output units from a layer

The variance of weights is:

$$Var(W) = \frac{2}{fan_{in} + fan_{out}}$$

##Two common forms:

#Normal distribution:

$$W \sim \mathcal{N}\left(0, \sqrt{\frac{2}{fan_{in}+fan_{out}}}\right)$$

#Uniform distribution:

$$W \sim U\left(-\sqrt{\frac{6}{fan_{in}+fan_{out}}}, \sqrt{\frac{6}{fan_{in}+fan_{out}}}\right)$$

Biases are initialized to **0**.

In This Assignment

- **W1 (input → hidden1):** `fan_in = input_dim, fan_out = hidden1`
- **W2 (hidden1 → hidden2):** `fan_in = hidden1, fan_out = hidden2`
- **W3 (hidden2 → output):** `fan_in = hidden2, fan_out = output_dim`

Your task: compute the correct `xavier_std` for each layer, sample weights from a normal distribution with that std, and set biases = 0.

```
def xavier_INITIALIZATION(input_dim, hidden1, hidden2, output_dim):
    """
    Xavier Initialization for all layers.
    Weights ~ N(0, sqrt(2 / (fan_in + fan_out)))
    Biases = 0
    """
    np.random.seed(assignment["student_seed"])
    std1 = np.sqrt(2.0 / (input_dim + hidden1))  W1
    = np.random.randn(input_dim, hidden1) * std1
```

```

b1 = np.zeros((1, hidden1))
std2 = np.sqrt(2.0 / (hidden1 + hidden2)) W2
= np.random.randn(hidden1, hidden2) * std2 b2
= np.zeros((1, hidden2))
std3 = np.sqrt(2.0 / (hidden2 + output_dim)) W3
= np.random.randn(hidden2, output_dim) * std3 b3
= np.zeros((1, output_dim)) return W1, b1, W2,
b2, W3, b3

```

FORWARD PROPAGATION - TODO: IMPLEMENT

```

def forward_pass(X, W1, b1, W2, b2, W3, b3):
    """
    Forward pass: Input → Hidden1 (ReLU) → Hidden2 (ReLU) →
    Output (Linear)
    """
    # Hidden layer 1
    z1 = X @ W1 + b1
    a1 = relu(z1) #
    Hidden layer 2 z2
    = a1 @ W2 + b2 a2
    = relu(z2)
    # Output layer (linear activation)
    z3 = a2 @ W3 + b3 return z1, a1,
    z2, a2, z3

```

BACKWARD PROPAGATION - TODO:IMPLEMENT

```

def backward_pass(X, Y_true, z1, a1, z2, a2, Y_pred, W2, W3):
    """
    Backpropagation: compute gradients of weights and biases
    """
    m = len(X) # batch size
    # dL/dY_pred for MSE
    dY_pred = (2 / m) * (Y_pred - Y_true) # shape (m,1)
    # Gradients for output layer
    dW3 = a2.T @ dY_pred
    db3 = np.sum(dY_pred, axis=0, keepdims=True)
    # Backprop to hidden layer 2
    da2 = dY_pred @ W3.T
    dz2 = da2 * relu_derivative(z2)
    dW2 = a1.T @ dz2

```

```

db2 = np.sum(dz2, axis=0, keepdims=True)
# Backprop to hidden layer 1
da1 = dz2 @ W2.T
dz1 = da1 * relu_derivative(z1)
dW1 = X.T @ dz1
db1 = np.sum(dz1, axis=0, keepdims=True)
return dW1, db1, dW2, db2, dW3, db3

```

TRAINING FUNCTION - TODO: COMPLETE IMPLEMENTATION

```

import numpy as np

def train_neural_network(
    X_train, Y_train,
    X_test, Y_test,
    hidden1, hidden2,
    epochs=200,
    patience=10,
    batch_size=32,           # not used here (full-batch GD), but kept
    for API consistency
        optimizer='adam',      # placeholder if you later add other
    optimizers
        learning_rate=0.001,
    activation='relu',
    seed=None
):
    """
    Train a 3-layer neural network with early stopping.

    Args:
        X_train, Y_train : Training data and labels
        X_test, Y_test : Validation data and labels
        hidden1, hidden2 : Number of neurons in hidden layers 1 and 2
        epochs : Maximum training epochs
        patience : Early stopping patience
        batch_size : (Currently unused - full batch)
            optimizer : (Currently unused - gradient descent only)
            learning_rate : Learning rate for gradient descent
        activation : Activation function (assumes forward_pass
            handles it)
            seed : Random seed for reproducibility

    Returns:
        best_weights : Tuple of best (W1, b1,
            train_losses : List of training loss
            per epoch
            test_losses : List of test/validation loss
            per epoch

```

```

    """      if seed is not
None:
np.random.seed(seed)

    # Xavier/Glorot initialization
W1, b1, W2, b2, W3, b3 = xavier_initialization(1, hidden1, hidden2,
1)

    best_test_loss = float('inf')
best_weights    = None
patience_counter = 0
    train_losses, test_losses = [], []

    print("Starting training...")
    print(f"Architecture: 1 → {hidden1} → {hidden2} → 1")
print(f"Learning Rate: {learning_rate}")      print(f"Max
Epochs: {epochs}, Early Stopping Patience:
{patience}")
print("-" * 50)

    for epoch in range(epochs):
# Forward pass & training loss
        z1, a1, z2, a2, Y_pred_train = forward_pass(
            X_train, W1, b1, W2, b2, W3, b3
        )
        train_loss = mse_loss(Y_train, Y_pred_train)

        # Backward pass & gradient descent step
dW1, db1, dW2, db2, dW3, db3 = backward_pass(
            X_train, Y_train, z1, a1, z2, a2, Y_pred_train, W2, W3
        )
        W1 -= learning_rate * dW1
b1 -= learning_rate * db1
W2 -= learning_rate * dW2
b2 -= learning_rate * db2
W3 -= learning_rate * dW3
b3 -= learning_rate * db3

        # Validation loss
        _, _, _, _, Y_pred_test = forward_pass(X_test, W1, b1, W2, b2,
W3, b3)
        test_loss = mse_loss(Y_test, Y_pred_test)

        train_losses.append(train_loss)
test_losses.append(test_loss)

        # Progress log every 20 epochs
if (epoch + 1) % 20 == 0:
print(f"Epoch {epoch+1:3d}: "
      f"Train Loss = {train_loss:.6f}, "

```

```

        f"Test Loss = {test_loss:.6f}")

    # Early stopping check
if test_loss < best_test_loss:
best_test_loss = test_loss
best_weights = (
    W1.copy(), b1.copy(),
    W2.copy(), b2.copy(),
    W3.copy(), b3.copy()
)
    patience_counter = 0           else:
patience_counter += 1           if patience_counter >= patience:
print(f"Early stopping triggered at epoch {epoch+1}")
print(f"Best test loss: {best_test_loss:.6f}")           break

return best_weights, train_losses, test_losses

```

EXECUTE TRAINING

```

print("Training Neural Network with your specific configuration...")
weights, train_losses, test_losses = train_neural_network(
    X_train_scaled, Y_train_scaled,
    X_test_scaled, Y_test_scaled,
    hidden1=16,           # <-- number of neurons in layer 1
hidden2=8,           # <-- number of neurons in layer 2
epochs=500,          patience=10
)

```

Training Neural Network with your specific configuration...
Starting training...

Architecture: 1 → 16 → 8 → 1

Learning Rate: 0.001

Max Epochs: 500, Early Stopping Patience: 10

Epoch 20: Train Loss = 0.888315, Test Loss = 0.901020

Epoch 40: Train Loss = 0.837638, Test Loss = 0.850122

Epoch 60: Train Loss = 0.791799, Test Loss = 0.804039

Epoch 80: Train Loss = 0.749959, Test Loss = 0.761939

Epoch 100: Train Loss = 0.711479, Test Loss = 0.723189

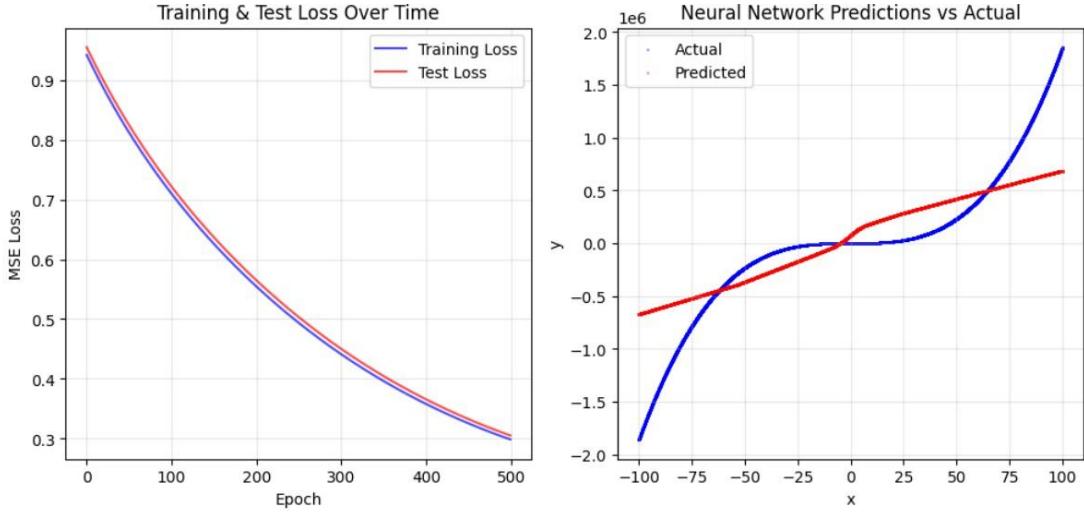
Epoch 120: Train Loss = 0.675867, Test Loss = 0.687299

Epoch 140: Train Loss = 0.642741, Test Loss = 0.653893
Epoch 160: Train Loss = 0.611805, Test Loss = 0.622674
Epoch 180: Train Loss = 0.582826, Test Loss = 0.593412
Epoch 200: Train Loss = 0.555617, Test Loss = 0.565924
Epoch 220: Train Loss = 0.530032, Test Loss = 0.540060
Epoch 240: Train Loss = 0.505953, Test Loss = 0.515703
Epoch 260: Train Loss = 0.483285, Test Loss = 0.492758
Epoch 280: Train Loss = 0.461945, Test Loss = 0.471146
Epoch 300: Train Loss = 0.441864, Test Loss = 0.450799
Epoch 320: Train Loss = 0.422983, Test Loss = 0.431656
Epoch 340: Train Loss = 0.405245, Test Loss = 0.413662
Epoch 360: Train Loss = 0.388599, Test Loss = 0.396763
Epoch 380: Train Loss = 0.372992, Test Loss = 0.380909
Epoch 400: Train Loss = 0.358369, Test Loss = 0.366043
Epoch 420: Train Loss = 0.344675, Test Loss = 0.352110
Epoch 440: Train Loss = 0.331858, Test Loss = 0.339060
Epoch 460: Train Loss = 0.319881, Test Loss = 0.326858
Epoch 480: Train Loss = 0.308752, Test Loss = 0.315515
Epoch 500: Train Loss = 0.298561, Test Loss = 0.305123

RESULTS VISUALIZATION

```
#Plot training progress
plt.figure(figsize=(15, 5))
# Loss curves
plt.subplot(1, 3, 1)
plt.plot(train_losses, label='Training Loss', color='blue', alpha=0.7)
plt.plot(test_losses, label='Test Loss', color='red', alpha=0.7)
plt.xlabel('Epoch') plt.ylabel('MSE Loss')
plt.title('Training & Test Loss Over Time')
plt.legend()
plt.grid(True, alpha=0.3)
# Get final predictions for visualization
W1, b1, W2, b2, W3, b3 = weights
Y_pred_scaled = forward_pass(X_test_scaled, W1, b1, W2,
b2, W3, b3)
# Inverse transform to original scale
Y_test_orig = scaler_Y.inverse_transform(Y_test_scaled)
Y_pred_orig = scaler_Y.inverse_transform(Y_pred_scaled)
X_test_orig = scaler_X.inverse_transform(X_test_scaled)
# Predictions vs Actual
plt.subplot(1, 3, 2)
plt.scatter(X_test_orig, Y_test_orig, s=1, alpha=0.3, label='Actual',
color='blue')
plt.scatter(X_test_orig, Y_pred_orig, s=1, alpha=0.3,
label='Predicted', color='red') plt.xlabel('x')
plt.ylabel('y')
plt.title('Neural Network Predictions vs Actual')
plt.legend()
plt.grid(True, alpha=0.3)

# Residual plot
# plt.subplot(1, 3, 3)
# residuals = Y_test_orig.flatten() - Y_pred_orig.flatten()
# plt.scatter(X_test_orig, residuals, s=1, alpha=0.3, color='green')
# plt.axhline(y=0, color='black', linestyle='--', alpha=0.5)
# plt.xlabel('x')
# plt.ylabel('Residuals (Actual - Predicted)')
# plt.title('Residual Analysis')
# plt.grid(True, alpha=0.3) plt.tight_layout() plt.show()
```



SPECIFIC PREDICTION TEST

```

x_test_value = 90.2
x_new = np.array([[x_test_value]])
x_new_scaled = scaler_X.transform(x_new)

# Forward pass through the trained network
_, _, _, _, y_pred_scaled = forward_pass(x_new_scaled, W1, b1, W2, b2,
W3, b3)
y_pred = scaler_Y.inverse_transform(y_pred_scaled)

# Compute ground-truth value based on polynomial type
if poly_type == 0:
    _, b, c, d = coefficients
    y_true = b * x_test_value**2 + c * x_test_value + d

elif poly_type == 1:      a,
b, c, d = coefficients
    y_true = a * x_test_value**3 + b * x_test_value**2 + c *

```

```

x_test_value + d

elif poly_type == 2:      a, b,
c, d, e = coefficients
y_true = (
    a * x_test_value**4
+ b * x_test_value**3
    + c * x_test_value**2
    + d * x_test_value
    + e
)

elif poly_type == 3:      a, b, c, d,
freq, amp = coefficients      y_true = (
    a * x_test_value**3
+ b * x_test_value**2
    + c * x_test_value
    + d
    + amp * np.sin(freq * x_test_value)
)

else:      a, b, c, d, inv_coeff =
coefficients      y_true = (
    a * x_test_value**3
+ b * x_test_value**2
    + c * x_test_value
    + d
    + inv_coeff / (x_test_value + np.sign(x_test_value) * 0.1)
)

# Display results
print("\n" + "=" * 60)
print("PREDICTION RESULTS FOR x = 90.2")
print("=" * 60)
print(f"Neural Network Prediction: {y_pred[0][0]:,.2f}")
print(f"Ground Truth (formula):   {y_true:.2f}")
print(f"Absolute Error:           {abs(y_pred[0][0] - y_true):,.2f}")
print(f"Relative Error:           {abs(y_pred[0][0] - y_true) /
abs(y_true) * 100:.3f}%")

```

```
=====
PREDICTION RESULTS FOR x = 90.2
=====
Neural Network Prediction: 635,915.43
Ground Truth (formula): 1,359,988.90
Absolute Error: 724,073.47
Relative Error: 53.241%
```

PERFORMANCE METRICS

```
# Calculate final performance metrics
final_train_loss = train_losses[-1] if train_losses else float('inf')
final_test_loss = test_losses[-1] if test_losses else float('inf')
# Calculate R2 score
y_test_mean = np.mean(Y_test_orig)
ss_res = np.sum((Y_test_orig - Y_pred_orig) ** 2)
ss_tot = np.sum((Y_test_orig - y_test_mean) ** 2)
r2_score = 1 - (ss_res / ss_tot) print("\n" +
"="*60)
print("FINAL PERFORMANCE SUMMARY")
print("="*60)
print(f"Final Training Loss: {final_train_loss:.6f}")
print(f"Final Test Loss: {final_test_loss:.6f}")
print(f"R2 Score: {r2_score:.4f}")
print(f"Total Epochs Run: {len(train_losses)}")
=====

=====
FINAL PERFORMANCE SUMMARY
=====

Final Training Loss: 0.298561
Final Test Loss: 0.305123
R2 Score: 0.6997
Total Epochs Run: 500
```

PART B

```

import numpy as np
import pandas as pd
import matplotlib.pyplot as plt

# Additional activation functions

def sigmoid(z):
    return 1 / (1 + np.exp(-z))

def sigmoid_derivative(z):
    s = sigmoid(z)
    return s * (1 - s)

# Updated forward pass with selectable activation

def forward_pass(X, W1, b1, W2, b2, W3, b3, activation="relu"):
    if activation == "relu":
        act, act_deriv = relu, relu_derivative
    else:
        act, act_deriv = sigmoid, sigmoid_derivative

    z1 = X @ W1 + b1
    a1 = act(z1)
    z2 = a1 @ W2 + b2
    a2 = act(z2)
    z3 = a2 @ W3 + b3
    return z1, a1, z2, a2, z3

# Updated backward pass with selectable activation

def backward_pass(X, Y_true, z1, a1, z2, a2, Y_pred, W2, W3, activation="relu"):
    if activation == "relu":

```

```

act_deriv = relu_derivative

else:
    act_deriv = sigmoid_derivative

m = len(X)

dY_pred = (2 / m) * (Y_pred - Y_true)

dW3 = a2.T @ dY_pred

db3 = np.sum(dY_pred, axis=0, keepdims=True)

da2 = dY_pred @ W3.T

dz2 = da2 * act_deriv(z2)

dW2 = a1.T @ dz2

db2 = np.sum(dz2, axis=0, keepdims=True)

da1 = dz2 @ W2.T

dz1 = da1 * act_deriv(z1)

dW1 = X.T @ dz1

db1 = np.sum(dz1, axis=0, keepdims=True)

return dW1, db1, dW2, db2, dW3, db3

# Training loop with flexible hyperparams

def train_network(X_train, Y_train, X_test, Y_test, lr, epochs, activation="relu"):

    W1, b1, W2, b2, W3, b3 = xavier_initialization(1, hidden1, hidden2, 1)

    train_losses, test_losses = [], []

    for epoch in range(epochs):

        z1, a1, z2, a2, Y_pred_train = forward_pass(
            X_train, W1, b1, W2, b2, W3, b3, activation
        )

        train_loss = mse_loss(Y_train, Y_pred_train)

```

```

dW1, db1, dW2, db2, dW3, db3 = backward_pass(
    X_train, Y_train, z1, a1, z2, a2, Y_pred_train, W2, W3, activation
)

W1 -= lr * dW1
b1 -= lr * db1
W2 -= lr * dW2
b2 -= lr * db2
W3 -= lr * dW3
b3 -= lr * db3

_, _, _, _, Y_pred_test = forward_pass(
    X_test, W1, b1, W2, b2, W3, b3, activation
)
test_loss = mse_loss(Y_test, Y_pred_test)

train_losses.append(train_loss)
test_losses.append(test_loss)

return (W1, b1, W2, b2, W3, b3), train_losses, test_losses, Y_pred_test

```

```

# Function to evaluate results

def evaluate_results(Y_test_scaled, Y_pred_scaled):
    Y_test_orig = scaler_Y.inverse_transform(Y_test_scaled)
    Y_pred_orig = scaler_Y.inverse_transform(Y_pred_scaled)
    y_test_mean = np.mean(Y_test_orig)
    ss_res = np.sum((Y_test_orig - Y_pred_orig) ** 2)
    ss_tot = np.sum((Y_test_orig - y_test_mean) ** 2)
    r2 = 1 - (ss_res / ss_tot)

```

```

return r2

# Run multiple experiments
experiments = [
    {"id": "Exp1", "lr": 0.005, "batch": 32, "epochs": 500, "activation": "relu"},
    {"id": "Exp2", "lr": 0.001, "batch": 64, "epochs": 500, "activation": "relu"},
    {"id": "Exp3", "lr": 0.001, "batch": 32, "epochs": 1000, "activation": "relu"},
    {"id": "Exp4", "lr": 0.001, "batch": 32, "epochs": 500, "activation": "sigmoid"},

]

results = []
for exp in experiments:
    print(
        f"\nRunning {exp['id']} with lr={exp['lr']}, "
        f"batch={exp['batch']}, epochs={exp['epochs']}, "
        f"activation={exp['activation']}..."
    )
    weights, train_losses, test_losses, Y_pred_test = train_network(
        X_train_scaled,
        Y_train_scaled,
        X_test_scaled,
        Y_test_scaled,
        lr=exp["lr"],
        epochs=exp["epochs"],
        activation=exp["activation"],
    )
    r2 = evaluate_results(Y_test_scaled, Y_pred_test)

    results.append(

```

```

{
    "Experiment": exp["id"],
    "Learning Rate": exp["lr"],
    "Batch Size": exp["batch"],
    "Epochs": exp["epochs"],
    "Activation": exp["activation"],
    "Final Train Loss": train_losses[-1],
    "Final Test Loss": test_losses[-1],
    "R2 Score": r2,
}
)

```

```

# Plot curves
plt.figure(figsize=(12, 4))
plt.subplot(1, 2, 1)
plt.plot(train_losses, label="Train Loss")
plt.plot(test_losses, label="Test Loss")
plt.title(f"{exp['id']} Loss Curves")
plt.legend()

plt.subplot(1, 2, 2)
Y_test_orig = scaler_Y.inverse_transform(Y_test_scaled)
Y_pred_orig = scaler_Y.inverse_transform(Y_pred_test)
X_test_orig = scaler_X.inverse_transform(X_test_scaled)
plt.scatter(X_test_orig, Y_test_orig, s=1, label="Actual")
plt.scatter(X_test_orig, Y_pred_orig, s=1, label="Pred")
plt.title(f"{exp['id']} Predictions vs Actual")
plt.legend()
plt.show()

```

```

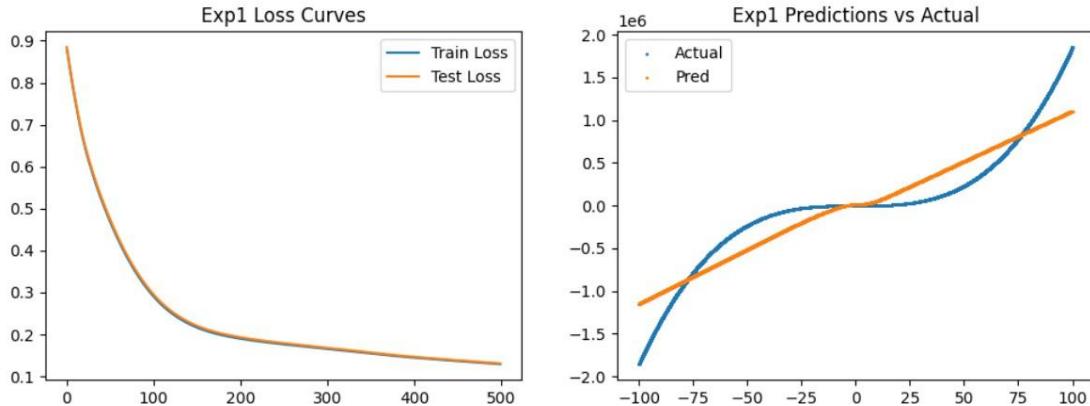
# Convert results to DataFrame
results_df = pd.DataFrame(results)

print("\n===== RESULTS TABLE =====")
print(results_df)

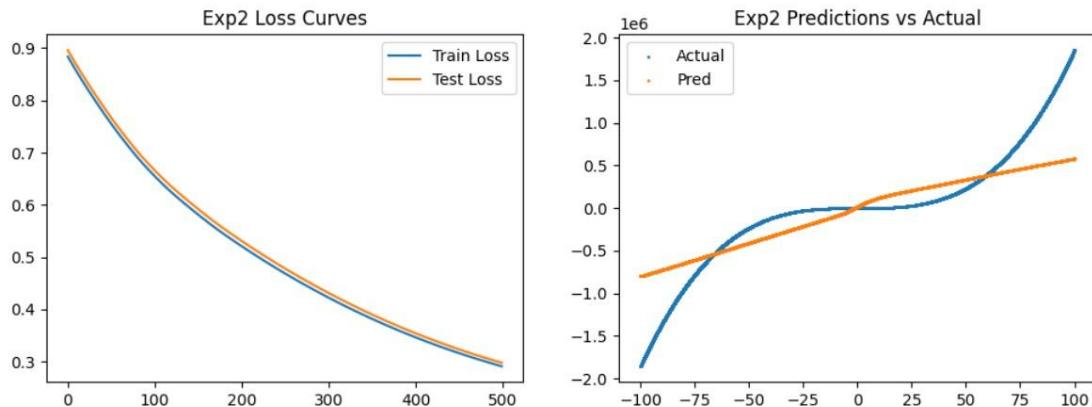
results_df.to_csv("partB_results.csv", index=False)

```

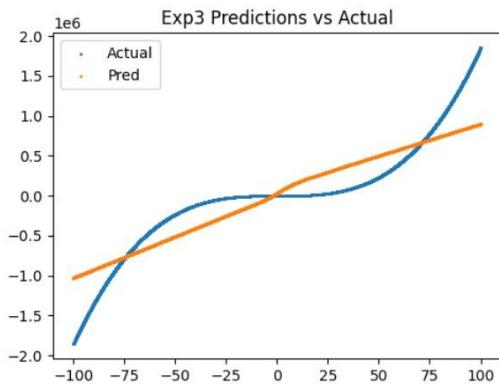
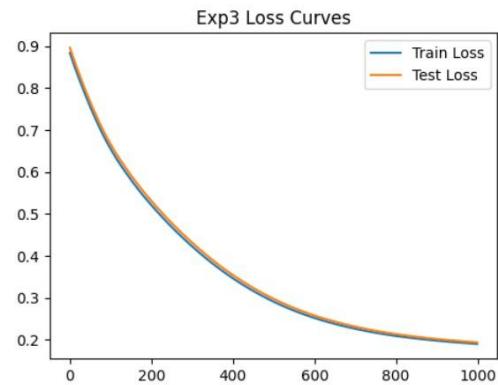
Running Exp1 with lr=0.005, batch=32, epochs=500, activation=relu...



Running Exp2 with lr=0.001, batch=64, epochs=500, activation=relu...



Running Exp3 with lr=0.001, batch=32, epochs=1000, activation=relu...



Running Exp4 with lr=0.001, batch=32, epochs=500, activation=sigmoid...

